

# Introduction to Callbacks

# Callbacks

- Provides some functionality at various stages of training
- Subclasses `tf.keras.callbacks.Callback`
- Useful in understanding a model's state during training
  - internal states
  - statistics e.g., losses and metrics

# Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None  
  
    def on_epoch_begin(self, epoch, logs=None):  
        """Called at the beginning of an epoch during training."""  
  
    def on_epoch_end(self, epoch, logs=None):  
        """Called at the end of an epoch during training."""
```

# Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None
```

```
def on_epoch_begin(self, epoch, logs=None):  
    """Called at the beginning of an epoch during training."""
```

```
def on_epoch_end(self, epoch, logs=None):  
    """Called at the end of an epoch during training."""
```

you can override these two when you want to customize the function

# Training specific methods

```
class Callback(object):  
    def __init__(self):  
        self.validation_data = None  
        self.model = None  
  
    def on_epoch_begin(self, epoch, logs=None):  
        """Called at the beginning of an epoch during training."""  
  
    def on_epoch_end(self, epoch, logs=None):  
        """Called at the end of an epoch during training."""
```

# Common methods for training/testing/predicting

```
class Callback(object):  
    ...  
    def on_(train|test|predict)_begin(self, logs=None):  
        """Called at the begin of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_end(self, logs=None):  
        """Called at the end of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_batch_begin(self, batch, logs=None):  
        """Called right before processing a batch during training/testing/predicting."""  
  
    def on_(train|test|predict)_batch_end(self, batch, logs=None):  
        """Called at the end of training/testing/predicting a batch."""
```

we can specify  
these callbacks at  
specific level  
during training  
and  
testing  
predicting

# Common methods for training/testing/predicting

```
class Callback(object):
```

```
...
```

```
def on_(train|test|predict)_begin(self, logs=None):
```

```
    """Called at the begin of fit/evaluate/predict."""
```

```
def on_(train|test|predict)_end(self, logs=None):
```

```
    """Called at the end of fit/evaluate/predict."""
```

```
def on_(train|test|predict)_batch_begin(self, batch, logs=None):
```

```
    """Called right before processing a batch during training/testing/predicting."""
```

```
def on_(train|test|predict)_batch_end(self, batch, logs=None):
```

```
    """Called at the end of training/testing/predicting a batch."""
```

# Common methods for training/testing/predicting

```
class Callback(object):  
    ...  
    def on_(train|test|predict)_begin(self, logs=None):  
        """Called at the begin of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_end(self, logs=None):  
        """Called at the end of fit/evaluate/predict."""  
  
    def on_(train|test|predict)_batch_begin(self, batch, logs=None):  
        """Called right before processing a batch during training/testing/predicting."""  
  
    def on_(train|test|predict)_batch_end(self, batch, logs=None):  
        """Called at the end of training/testing/predicting a batch."""
```



# Where can you use them?

Model methods that take callbacks

- `fit(..., callbacks=[...])`
- `fit_generator(..., callbacks=[...])`
- `evaluate(..., callbacks=[...])`
- `evaluate_generator(..., callbacks=[...])`
- `predict(..., callbacks=[...])`
- `predict_generator(..., callbacks=[...])`

# TensorBoard Callback

- Visualize machine learning experiments
- Track metrics (e.g., loss, accuracy)
- View the model graph

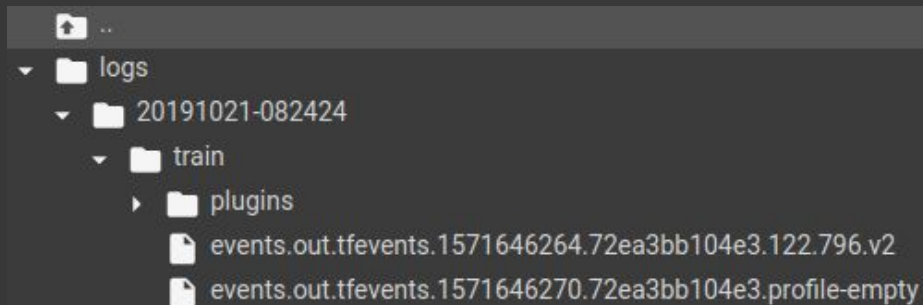
*can be used  
across these area*

```
TensorBoard(log_dir='./logs', update_freq='epoch', **kwargs)
```


<https://www.tensorflow.org/tensorboard>

```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))  
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)  
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```


it is initialized at the  
TensorBoard callback.



```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))  
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)  
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```



```
log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))  
tensorboard = tf.keras.callbacks.TensorBoard(log_dir=log_dir)  
model.fit(train_batches, epochs=10, callbacks=[tensorboard])
```

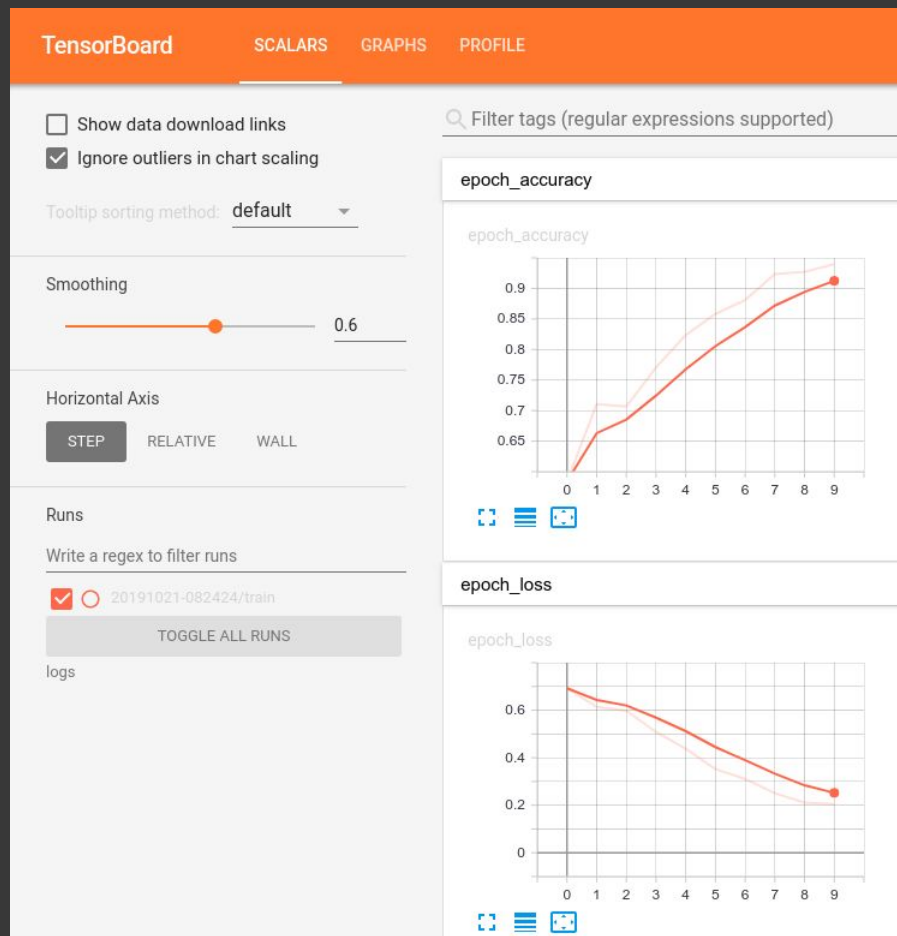


```
# Load the extension
```

```
%load_ext tensorboard
```

```
# Run TensorBoard
```

```
%tensorboard --logdir logs
```



# Model Checkpoints

# ModelCheckpoint

- Saves the model every so often
- Choose to save only the best checkpoints / weights

```
ModelCheckpoint(filepath, monitor='val_loss', mode='auto',  
                save_best_only=False, save_weights_only=False,  
                verbose=0, save_freq=1, **kwargs)
```



```
model.fit(train_batches, epochs=5, validation_data=validation_batches,  
          callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val\_loss: 0.6608 - val\_accuracy: 0.8667

Epoch 3/5

Epoch 00003: saving model to model.h5

33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val\_loss: 0.6217 - val\_accuracy: 0.9417

Epoch 4/5

Epoch 00004: saving model to model.h5

33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val\_loss: 0.5591 - val\_accuracy: 0.8667

Epoch 5/5

Epoch 00005: saving model to model.h5

33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val\_loss: 0.4485 - val\_accuracy: 0.8583

<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,  
         callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val\_loss: 0.6608 - val\_accuracy: 0.8667

Epoch 3/5

Epoch 00003: saving model to model.h5

33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val\_loss: 0.6217 - val\_accuracy: 0.9417

Epoch 4/5

Epoch 00004: saving model to model.h5

33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val\_loss: 0.5591 - val\_accuracy: 0.8667

Epoch 5/5

Epoch 00005: saving model to model.h5

33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val\_loss: 0.4485 - val\_accuracy: 0.8583

<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches,  
         callbacks=[ModelCheckpoint('model.h5', verbose=1)])
```

Epoch 1/5

```
Epoch 00001: saving model to model.h5  
33/33 - 7s - loss: 0.6879 - accuracy: 0.6702 - val_loss: 0.0000e+00 - val_accuracy: 0.0000e+00  
Epoch 2/5
```

```
Epoch 00002: saving model to model.h5  
33/33 - 6s - loss: 0.6721 - accuracy: 0.8447 - val_loss: 0.6608 - val_accuracy: 0.8667  
Epoch 3/5
```

```
Epoch 00003: saving model to model.h5  
33/33 - 6s - loss: 0.6435 - accuracy: 0.8840 - val_loss: 0.6217 - val_accuracy: 0.9417  
Epoch 4/5
```

```
Epoch 00004: saving model to model.h5  
33/33 - 6s - loss: 0.5920 - accuracy: 0.8849 - val_loss: 0.5591 - val_accuracy: 0.8667  
Epoch 5/5
```

```
Epoch 00005: saving model to model.h5  
33/33 - 6s - loss: 0.5047 - accuracy: 0.9089 - val_loss: 0.4485 - val_accuracy: 0.8583  
<tensorflow.python.keras.callbacks.History at 0x7f09ccef97f0>
```

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,  
          callbacks=[ModelCheckpoint('model.h5', save_weights_only=True, verbose=1)])
```

Epoch 1/2

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6493 - accuracy: 0.6184 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.5684 - accuracy: 0.7507 - val\_loss: 0.5183 - val\_accuracy: 0.7083

<tensorflow.python.keras.callbacks.History at 0x7f09cb5547f0>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,  
          callbacks=[ModelCheckpoint('model.h5', monitor='val_loss',  
                                     save_best_only=True, verbose=1)])
```

Epoch 1/5

Epoch 00001: val\_loss improved from inf to 0.65278, saving model to model.h5  
33/33 - 7s - loss: 0.6753 - accuracy: 0.5772 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00  
Epoch 2/5

Epoch 00002: val\_loss improved from 0.65278 to 0.62279, saving model to model.h5  
33/33 - 6s - loss: 0.6219 - accuracy: 0.7584 - val\_loss: 0.6228 - val\_accuracy: 0.5417  
Epoch 3/5

Epoch 00003: val\_loss improved from 0.62279 to 0.47633, saving model to model.h5  
33/33 - 6s - loss: 0.5448 - accuracy: 0.7977 - val\_loss: 0.4763 - val\_accuracy: 0.8750  
Epoch 4/5

Epoch 00004: val\_loss improved from 0.47633 to 0.44497, saving model to model.h5  
33/33 - 6s - loss: 0.4673 - accuracy: 0.8054 - val\_loss: 0.4450 - val\_accuracy: 0.8000  
Epoch 5/5

Epoch 00005: val\_loss improved from 0.44497 to 0.30997, saving model to model.h5  
33/33 - 6s - loss: 0.4030 - accuracy: 0.8677 - val\_loss: 0.3100 - val\_accuracy: 0.9000  
<tensorflow.python.keras.callbacks.History at 0x7f09cc9b7128>

```
model.fit(train_batches, epochs=5, validation_data=validation_batches, verbose=2,  
          callbacks=[ModelCheckpoint('model.h5', monitor='val_loss',  
                                     save_best_only=True, verbose=1)])
```

Epoch 1/5

Epoch 00001: val\_loss improved from inf to 0.65278, saving model to model.h5  
33/33 - 7s - loss: 0.6753 - accuracy: 0.5772 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00  
Epoch 2/5

Epoch 00002: val\_loss improved from 0.65278 to 0.62279, saving model to model.h5  
33/33 - 6s - loss: 0.6219 - accuracy: 0.7584 - val\_loss: 0.6228 - val\_accuracy: 0.5417  
Epoch 3/5

Epoch 00003: val\_loss improved from 0.62279 to 0.47633, saving model to model.h5  
33/33 - 6s - loss: 0.5448 - accuracy: 0.7977 - val\_loss: 0.4763 - val\_accuracy: 0.8750  
Epoch 4/5

Epoch 00004: val\_loss improved from 0.47633 to 0.44497, saving model to model.h5  
33/33 - 6s - loss: 0.4673 - accuracy: 0.8054 - val\_loss: 0.4450 - val\_accuracy: 0.8000  
Epoch 5/5

Epoch 00005: val\_loss improved from 0.44497 to 0.30997, saving model to model.h5  
33/33 - 6s - loss: 0.4030 - accuracy: 0.8677 - val\_loss: 0.3100 - val\_accuracy: 0.9000  
<tensorflow.python.keras.callbacks.History at 0x7f09cc9b7128>

```
model.fit(..., callbacks=[ModelCheckpoint('saved_model', ...)])
```

Epoch 1/2

Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6714 - accuracy: 0.5695 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6238 - accuracy: 0.6366 - val\_loss: 0.6459 - val\_accuracy: 0.5417

```
model.fit(..., callbacks=[ModelCheckpoint('model.h5', ...)])
```

Epoch 1/2

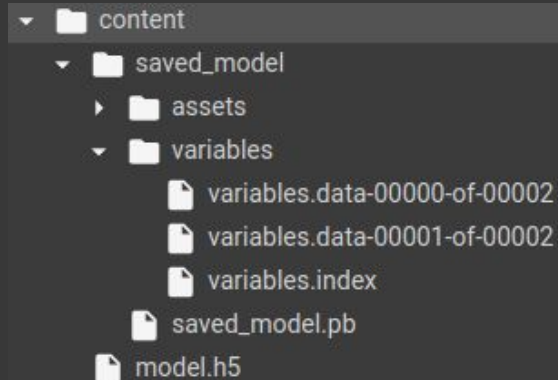
Epoch 00001: saving model to model.h5

33/33 - 7s - loss: 0.6714 - accuracy: 0.5695 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/2

Epoch 00002: saving model to model.h5

33/33 - 6s - loss: 0.6238 - accuracy: 0.6366 - val\_loss: 0.6459 - val\_accuracy: 0.5417



```
model.fit(..., callbacks=[ModelCheckpoint('weights.{epoch:02d}-{val_loss:.2f}.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to weights.01-0.63.h5

33/33 - 6s - loss: 0.6709 - accuracy: 0.6098 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to weights.02-0.60.h5

33/33 - 6s - loss: 0.6088 - accuracy: 0.7124 - val\_loss: 0.6046 - val\_accuracy: 0.5917

Epoch 3/5

Epoch 00003: saving model to weights.03-0.46.h5

33/33 - 6s - loss: 0.5354 - accuracy: 0.7613 - val\_loss: 0.4602 - val\_accuracy: 0.8500

Epoch 4/5

Epoch 00004: saving model to weights.04-0.38.h5

33/33 - 6s - loss: 0.4769 - accuracy: 0.7891 - val\_loss: 0.3848 - val\_accuracy: 0.9250

Epoch 5/5

Epoch 00005: saving model to weights.05-0.33.h5

33/33 - 6s - loss: 0.3961 - accuracy: 0.8600 - val\_loss: 0.3263 - val\_accuracy: 0.8667

content

weights.01-0.63.h5

weights.02-0.60.h5

weights.03-0.46.h5

weights.04-0.38.h5

weights.05-0.33.h5



```
model.fit(..., callbacks=[ModelCheckpoint('weights.{epoch:02d}-{val_loss:.2f}.h5', verbose=1)])
```

Epoch 1/5

Epoch 00001: saving model to weights.01-0.63.h5

33/33 - 6s - loss: 0.6709 - accuracy: 0.6098 - val\_loss: 0.0000e+00 - val\_accuracy: 0.0000e+00

Epoch 2/5

Epoch 00002: saving model to weights.02-0.60.h5

33/33 - 6s - loss: 0.6088 - accuracy: 0.7124 - val\_loss: 0.6046 - val\_accuracy: 0.5917

Epoch 3/5

Epoch 00003: saving model to weights.03-0.46.h5

33/33 - 6s - loss: 0.5354 - accuracy: 0.7613 - val\_loss: 0.4602 - val\_accuracy: 0.8500

Epoch 4/5

Epoch 00004: saving model to weights.04-0.38.h5

33/33 - 6s - loss: 0.4769 - accuracy: 0.7891 - val\_loss: 0.3848 - val\_accuracy: 0.9250

Epoch 5/5

Epoch 00005: saving model to weights.05-0.33.h5

33/33 - 6s - loss: 0.3961 - accuracy: 0.8600 - val\_loss: 0.3263 - val\_accuracy: 0.8667

content

weights.01-0.63.h5

weights.02-0.60.h5

weights.03-0.46.h5

weights.04-0.38.h5

weights.05-0.33.h5

# EarlyStopping

- Helps you keep track of a certain metric/loss and change training behavior accordingly
- Stops training when there's no improvement observed

```
EarlyStopping(monitor='val_loss', verbose=0,  
              min_delta=0, patience=0, mode='auto',  
              baseline=None, restore_best_weights=False,  
              **kwargs)
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(train_batches, epochs=50, validation_data=validation_batches,  
          callbacks=[EarlyStopping(patience=3, monitor='val_loss')])
```

```
Epoch 11/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2423 - accuracy: 0.9319 - val_loss: 0.1474 - val_accuracy: 0.9500  
Epoch 12/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1521 - accuracy: 0.9607 - val_loss: 0.1990 - val_accuracy: 0.9000  
Epoch 13/50  
33/33 [=====] - 6s 182ms/step - loss: 0.1571 - accuracy: 0.9511 - val_loss: 0.1176 - val_accuracy: 0.9500  
Epoch 14/50  
33/33 [=====] - 6s 186ms/step - loss: 0.1409 - accuracy: 0.9569 - val_loss: 0.1071 - val_accuracy: 0.9583  
Epoch 15/50  
33/33 [=====] - 6s 185ms/step - loss: 0.1450 - accuracy: 0.9626 - val_loss: 0.0953 - val_accuracy: 0.9583  
Epoch 16/50  
33/33 [=====] - 6s 184ms/step - loss: 0.2023 - accuracy: 0.9396 - val_loss: 0.1413 - val_accuracy: 0.9583  
Epoch 17/50  
33/33 [=====] - 6s 184ms/step - loss: 0.1443 - accuracy: 0.9655 - val_loss: 0.1771 - val_accuracy: 0.9167  
Epoch 18/50  
33/33 [=====] - 6s 183ms/step - loss: 0.1442 - accuracy: 0.9521 - val_loss: 0.1201 - val_accuracy: 0.9333  
Epoch 00018: early stopping
```

```
model.fit(...,  
          callbacks=[EarlyStopping(patience=3, restore_best_weights=True,  
                                   monitor='val_loss', verbose=1)])
```

```
Epoch 11/50  
33/33 - 6s - loss: 0.1380 - accuracy: 0.9616 - val_loss: 0.0968 - val_accuracy: 0.9750  
Epoch 12/50  
33/33 - 6s - loss: 0.1202 - accuracy: 0.9655 - val_loss: 0.0741 - val_accuracy: 0.9917  
Epoch 13/50  
33/33 - 6s - loss: 0.1716 - accuracy: 0.9434 - val_loss: 0.1083 - val_accuracy: 0.9750  
Epoch 14/50  
33/33 - 6s - loss: 0.1331 - accuracy: 0.9626 - val_loss: 0.0861 - val_accuracy: 0.9667  
Epoch 15/50  
Restoring model weights from the end of the best epoch.  
33/33 - 6s - loss: 0.1393 - accuracy: 0.9578 - val_loss: 0.0771 - val_accuracy: 0.9750  
Epoch 00015: early stopping
```

```
model.fit(...,  
           callbacks=[EarlyStopping(  
                                   patience=3,  
                                   min_delta=0.05,  
                                   baseline=0.8,  
                                   mode='min',  
                                   monitor='val_loss',  
                                   verbose=1  
                               )])
```



```
model.fit(..., callbacks=[CSVLogger('training.csv')])
```

epoch	accuracy	loss	val_accuracy	val_loss
0	0.574305	0.682536	0.775000	0.655427
1	0.760307	0.633610	0.675000	0.595201
2	0.758389	0.573186	0.850000	0.503174
3	0.835091	0.472031	0.808333	0.416691
4	0.854267	0.419491	0.916667	0.309128

## Multiple callbacks

```
model.fit(..., callbacks=[EarlyStopping(...),
model.evaluate(...      ModelCheckpoint(...),
model.predict(...      TensorBoard(...),
                        ...
                        ])
```

# Build a simple model

```
model = tf.keras.Sequential()
```

```
model.add(tf.keras.layers.Dense(units=1,  
                                activation='linear',  
                                input_dim=(784,)))
```

```
model.compile(optimizer=tf.keras.optimizers.RMSprop(lr=0.1),  
              loss='mean_squared_error', metrics=['mae'])
```

# How a custom callback looks

```
import datetime

class MyCustomCallback(tf.keras.callbacks.Callback):
    def on_train_batch_begin(self, batch, logs=None):
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))

    def on_train_batch_end(self, batch, logs=None):
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))
```

# How a custom callback looks

```
import datetime
```

```
class MyCustomCallback(tf.keras.callbacks.Callback):
```

```
    def on_train_batch_begin(self, batch, logs=None):
```

```
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.now().time()))
```

```
    def on_train_batch_end(self, batch, logs=None):
```

```
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.now().time()))
```

```
my_custom_callback = MyCustomCallback()
```

```
model.fit(x_train, y_train, batch_size=64, epochs=1, verbose=0,  
          callbacks=[my_custom_callback])
```

```
my_custom_callback = MyCustomCallback()
```

```
model.fit(x_train, y_train, batch_size=64, epochs=1, verbose=0,  
          callbacks=[my_custom_callback])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):
    def __init__(self, threshold):
        super(DetectOverfittingCallback, self).__init__()
        self.threshold = threshold

    def on_epoch_end(self, epoch, logs=None):
        ratio = logs["val_loss"] / logs["loss"]
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))

        if ratio > threshold:
            print("Stopping training...")
            self.model.stop_training = True

model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```



```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold  
  
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))  
  
        if ratio > threshold:  
            print("Stopping training...")  
            self.model.stop_training = True  
  
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):
```

```
    def __init__(self, threshold):
```

```
        super(DetectOverfittingCallback, self).__init__()
```

```
        self.threshold = threshold
```

```
    def on_epoch_end(self, epoch, logs=None):
```

```
        ratio = logs["val_loss"] / logs["loss"]
```

```
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))
```

```
        if ratio > threshold:
```

```
            print("Stopping training...")
```

```
            self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold
```

```
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))
```

```
    if ratio > threshold:  
        print("Stopping training...")  
        self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```



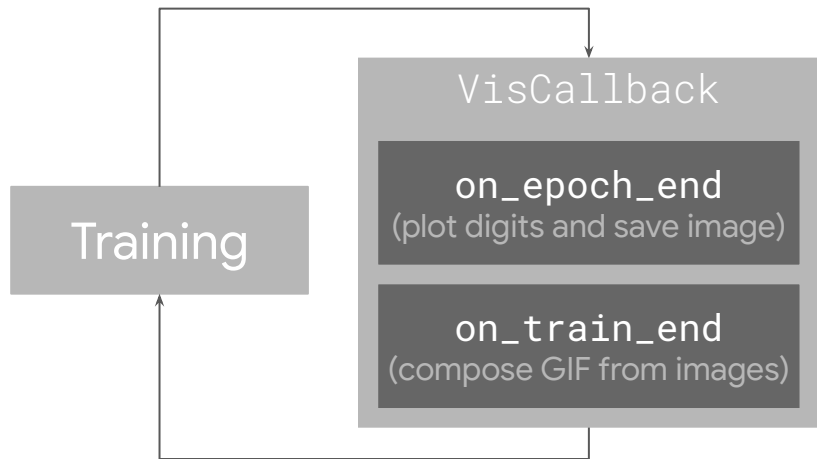
```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):
    def __init__(self, threshold):
        super(DetectOverfittingCallback, self).__init__()
        self.threshold = threshold

    def on_epoch_end(self, epoch, logs=None):
        ratio = logs["val_loss"] / logs["loss"]
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))

        if ratio > threshold:
            print("Stopping training...")
            self.model.stop_training = True
```

```
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```

```
class DetectOverfittingCallback(tf.keras.callbacks.Callback):  
    def __init__(self, threshold):  
        super(DetectOverfittingCallback, self).__init__()  
        self.threshold = threshold  
  
    def on_epoch_end(self, epoch, logs=None):  
        ratio = logs["val_loss"] / logs["loss"]  
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))  
  
        if ratio > threshold:  
            print("Stopping training...")  
            self.model.stop_training = True  
  
model.fit(..., callbacks=[DetectOverfittingCallback(threshold=1.3)])
```



```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                  n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```



```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                  n_samples=10):  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                 n_samples=10):  
  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                  n_samples=10):  
  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    def __init__(self, inputs, ground_truth, display_freq=10,  
                  n_samples=10):  
  
        self.inputs = inputs  
        self.ground_truth = ground_truth  
        self.images = []  
        self.display_freq = display_freq  
        self.n_samples = n_samples
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        # Randomly sample data  
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)  
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]  
        predictions = np.argmax(self.model.predict(X_test), axis=1)
```



```
class VisCallback(tf.keras.callbacks.Callback):
    ...
    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

        # Plot the digits
        display_digits(X_test, predictions, y_test, epoch, n=self.display_freq)
        ...
```

```
class VisCallback(tf.keras.callbacks.Callback):
```

```
...
```

```
def on_epoch_end(self, epoch, logs=None):
```

```
...
```

```
# Save the figure
```

```
buf = io.BytesIO()
```

```
plt.savefig(buf, format='png')
```

```
buf.seek(0)
```

```
image = Image.open(buf)
```

```
self.images.append(np.array(image))
```

```
# Display the digits every now and then
```

```
if epoch % self.display_freq == 0:
```

```
    plt.show()
```

```
class VisCallback(tf.keras.callbacks.Callback):  
    ...  
    def on_epoch_end(self, epoch, logs=None):  
        ...  
        # Save the figure  
        buf = io.BytesIO()  
        plt.savefig(buf, format='png')  
        buf.seek(0)  
        image = Image.open(buf)  
        self.images.append(np.array(image))  
  
        # Display the digits every now and then  
        if epoch % self.display_freq == 0:  
            plt.show()
```

```
import imageio
```

```
class VisCallback(tf.keras.callbacks.Callback):
```

```
...
```

```
def on_train_end(self, logs=None):
```

```
    imageio.mimsave('animation.gif', self.images, fps=1)
```

```
# Train the model
```

```
model.fit(..., callbacks=[VisCallback(x_test, y_test)])
```

```
import imageio
```

```
class VisCallback(tf.keras.callbacks.Callback):
```

```
    ...
```

```
    def on_train_end(self, logs=None):
```

```
        imageio.mimsave('animation.gif', self.images, fps=1)
```

```
# Train the model
```

```
model.fit(..., callbacks=[VisCallback(x_test, y_test)])
```

9 1 4 6 3 2 8 9 0 5